



```
!sudo apt-get install -y  
!sudo fc-cache -fv  
!rm ~/.cache/matplotlib -rf
```




```
%matplotlib inline  
import pandas as pd  
import numpy as nd  
import matplotlib  
import matplotlib.pyplot as plt  
import warnings  
warnings.filterwarnings('ignore')  
wine_data = pd.read_csv("/content/wine.csv", encoding="UTF-8")  
wine_data.head(10)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4	5
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4	5
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0	7
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	9.5	7
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5	5


```
wine_data.boxplot(column = "pH", by = "quality")
```

```
<Axes: title={'center': 'pH'}, xlabel='quality'>
```



```
plt.scatter(x=wine_data['fixed_acidity'], y=wine_data['density'])
```

```
<matplotlib.collections.PathCollection at 0x7b8458e9dd50>
```

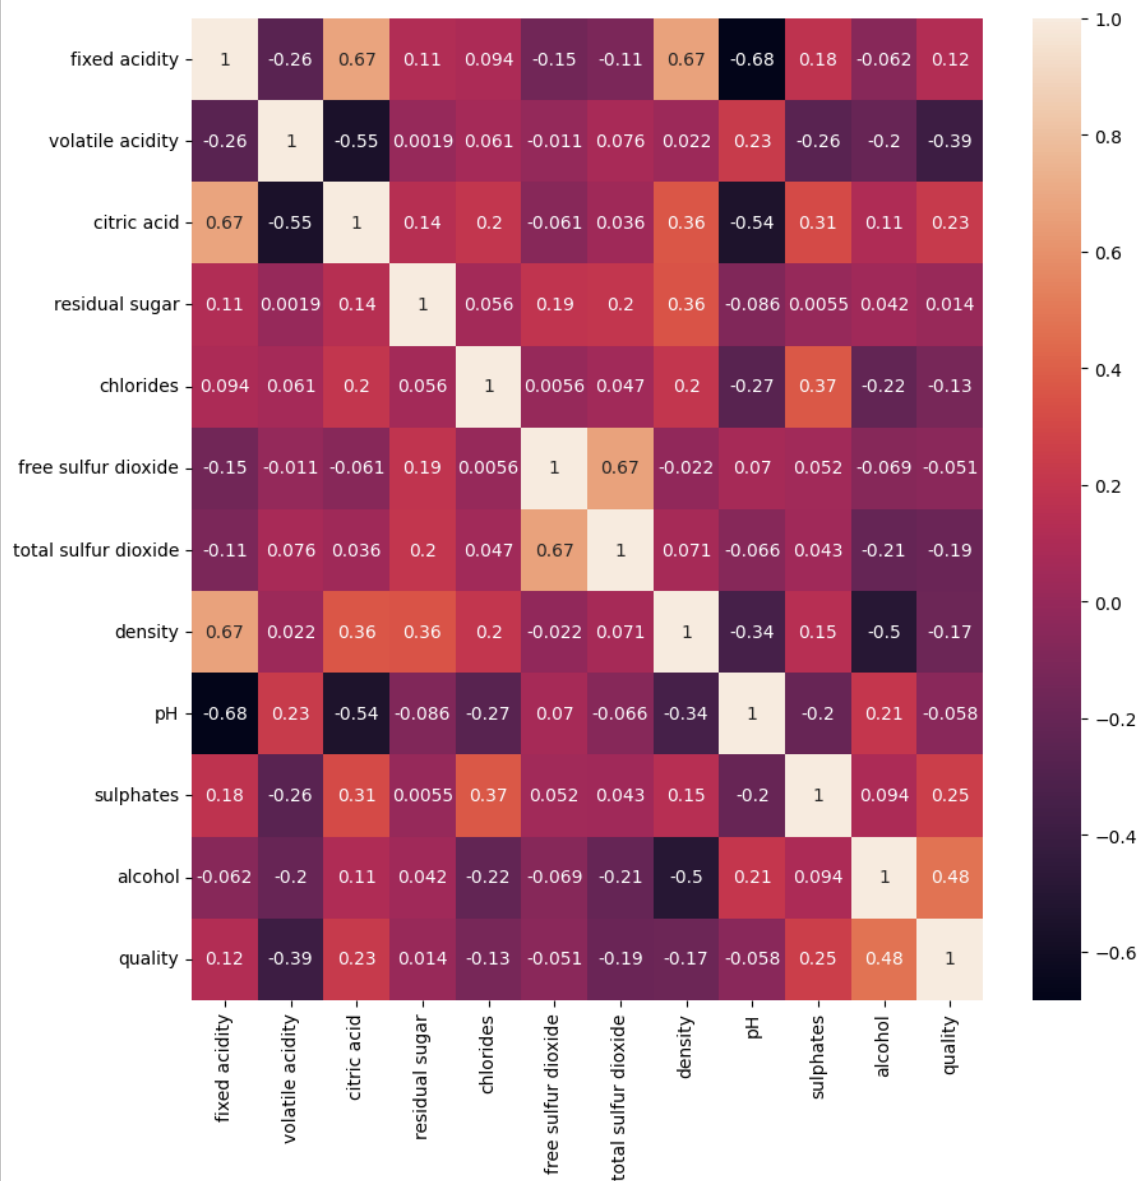


[5] wine\_data.corr()

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
fixed acidity	1.000000	-0.256131	0.671703	0.114777	0.093705	-0.153794	-0.113181	0.668047	-0.682978	0.183006	-0.061668	0.124052
volatile acidity	-0.256131	1.000000	-0.552496	0.001918	0.061298	-0.010504	0.076470	0.022026	0.234937	-0.260987	-0.202288	-0.390558
citric acid	0.671703	-0.552496	1.000000	0.143577	0.203823	-0.060978	0.035533	0.364947	-0.541904	0.312770	0.109903	0.226373
residual sugar	0.114777	0.001918	0.143577	1.000000	0.055610	0.187049	0.203028	0.355283	-0.085652	0.005527	0.042075	0.013732
chlorides	0.093705	0.061298	0.203823	0.055610	1.000000	0.005562	0.047400	0.200632	-0.265026	0.371260	-0.221141	-0.128907
free sulfur dioxide	-0.153794	-0.010504	-0.060978	0.187049	0.005562	1.000000	0.667666	-0.021946	0.070377	0.051658	-0.069408	-0.050656
total sulfur dioxide	-0.113181	0.076470	0.035533	0.203028	0.047400	0.667666	1.000000	0.071269	-0.066495	0.042947	-0.205654	-0.185100
density	0.668047	0.022026	0.364947	0.355283	0.200632	-0.021946	0.071269	1.000000	-0.341699	0.148506	-0.496180	-0.174919
pH	-0.682978	0.234937	-0.541904	-0.085652	-0.265026	0.070377	-0.066495	-0.341699	1.000000	-0.196648	0.205633	-0.057731
sulphates	0.183006	-0.260987	0.312770	0.005527	0.371260	0.051658	0.042947	0.148506	-0.196648	1.000000	0.093595	0.251397
alcohol	-0.061668	-0.202288	0.109903	0.042075	-0.221141	-0.069408	-0.205654	-0.496180	0.205633	0.093595	1.000000	0.476166
quality	0.124052	-0.390558	0.226373	0.013732	-0.128907	-0.050656	-0.185100	-0.174919	-0.057731	0.251397	0.476166	1.000000

```
[6] import seaborn as sns
plt.figure(figsize = (10,10))
sns.heatmap(data = wine_data.corr(), annot = True, color = 'red')
```

<Axes: >



```
[7] wine_data = pd.read_csv("/content/wine.csv", encoding="UTF-8")
print(wine_data.columns)
```

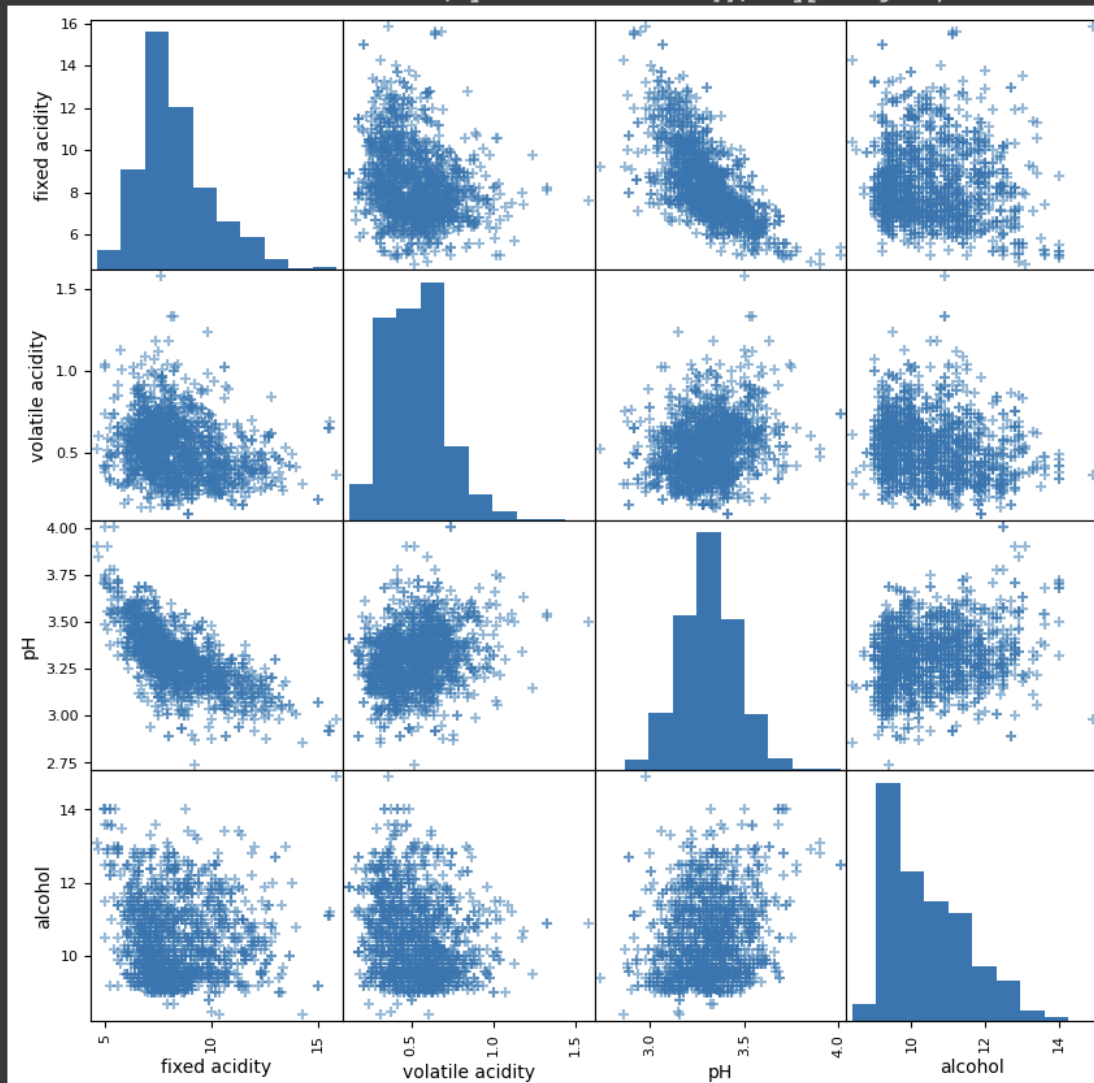
```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
       'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
       'pH', 'sulphates', 'alcohol', 'quality'],
      dtype='object')
```

```
[8] wine_info = wine_data[['fixed acidity', 'volatile acidity', 'pH', 'alcohol']]
wine_info.describe()
```

	fixed acidity	volatile acidity	pH	alcohol
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	3.311113	10.422983
std	1.741096	0.179060	0.154386	1.065668
min	4.600000	0.120000	2.740000	8.400000
25%	7.100000	0.390000	3.210000	9.500000
50%	7.900000	0.520000	3.310000	10.200000
75%	9.200000	0.640000	3.400000	11.100000
max	15.900000	1.580000	4.010000	14.900000

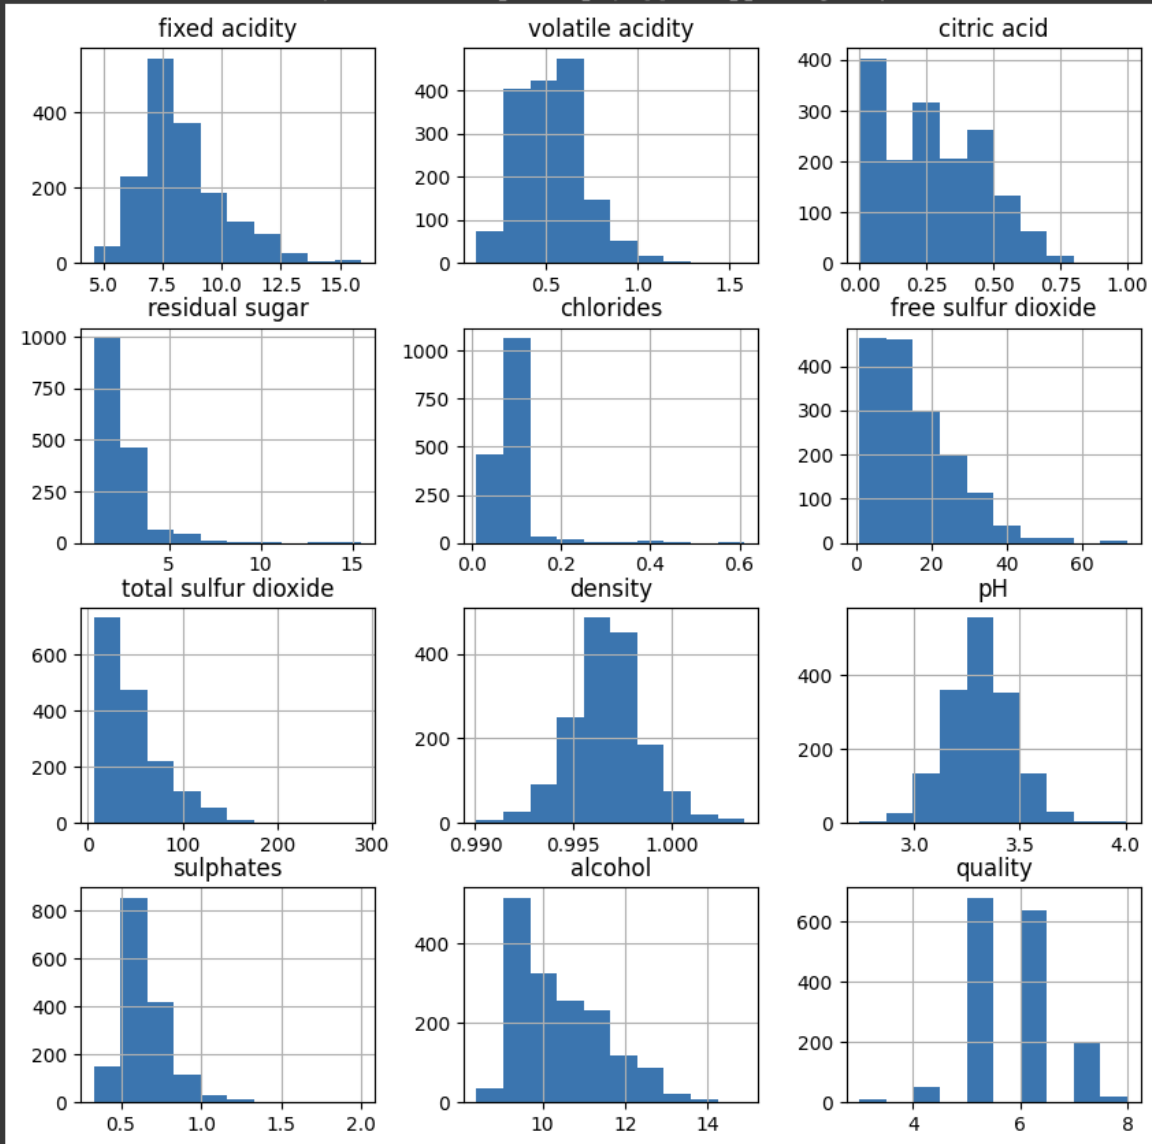
```
[9] pd.plotting.scatter_matrix(wine_info, marker = '+', figsize = (10,10))
```

```
<Axes: xlabel='volatile acidity', ylabel='pH'>,  
<Axes: xlabel='pH', ylabel='pH'>,  
<Axes: xlabel='alcohol', ylabel='pH'>],  
[<Axes: xlabel='fixed acidity', ylabel='alcohol'>,  
<Axes: xlabel='volatile acidity', ylabel='alcohol'>,  
<Axes: xlabel='pH', ylabel='alcohol'>,  
<Axes: xlabel='alcohol', ylabel='alcohol'>]], dtype=object)
```



```
[10] wine_data.hist(figsize=(10,10))
```

```
<Axes: title={'center': 'free sulfur dioxide'}>],  
[<Axes: title={'center': 'total sulfur dioxide'}>],  
<Axes: title={'center': 'density'}>],  
<Axes: title={'center': 'pH'}>],  
[<Axes: title={'center': 'sulphates'}>],  
<Axes: title={'center': 'alcohol'}>],  
<Axes: title={'center': 'quality'}>]], dtype=object)
```



```

[11] import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score, classification_report

[12] wine_data = pd.read_csv("/content/wine.csv", encoding="UTF-8")

[14] X = wine_data.drop(['quality'], axis=1)
      y = wine_data['quality'] # the target class

      # Split data into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

      # Initialize the Random Forest Classifier
      rf_clf = RandomForestClassifier(n_estimators=100, random_state=42)

      # Train the model
      rf_clf.fit(X_train, y_train)

      # Predict on the test data
      y_pred = rf_clf.predict(X_test)

      # Evaluate the model
      accuracy = accuracy_score(y_test, y_pred)
      print(f"Accuracy: {accuracy}")
      print(classification_report(y_test, y_pred))

```

```

Accuracy: 0.659375

```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	1
4	0.00	0.00	0.00	10
5	0.72	0.75	0.73	130
6	0.63	0.69	0.66	132
7	0.63	0.52	0.57	42
8	0.00	0.00	0.00	5
accuracy			0.66	320
macro avg	0.33	0.33	0.33	320
weighted avg	0.63	0.66	0.64	320



Seems that the support values suggested are imbalanced. Some classes have more values for accurate predictions and others dont. This time, I will use XGBoost.

```
[16] !pip install xgboost
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.10/dist-packages (2.0.2)  
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.23.5)  
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.11.3)
```

```
[18] import pandas as pd  
from xgboost import XGBClassifier  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import classification_report, accuracy_score  
from sklearn.preprocessing import LabelEncoder  
  
# Load data  
wine_data = pd.read_csv("./wine.csv", encoding="UTF-8")  
  
# Prepare data  
X = wine_data.drop('quality', axis=1)  
y = wine_data['quality']  
  
# Reindex classes to start from 0  
label_encoder = LabelEncoder()  
y_encoded = label_encoder.fit_transform(y)  
  
# Split data  
X_train, X_test, y_train_encoded, y_test_encoded = train_test_split(X, y_encoded, test_size=0.2, random_state=42)  
  
# Initialize XGBoost Classifier  
xgb_clf = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')  
  
# Train the model  
xgb_clf.fit(X_train, y_train_encoded)  
  
# Make predictions  
y_pred_encoded = xgb_clf.predict(X_test)  
y_pred = label_encoder.inverse_transform(y_pred_encoded)  
  
# Evaluation  
print("Accuracy:", accuracy_score(y_test_encoded, y_pred_encoded))  
print(classification_report(y_test_encoded, y_pred_encoded))
```

```
Accuracy: 0.696875  
      precision    recall  f1-score   support  
  
     0         0.00         0.00         0.00         1  
     1         0.00         0.00         0.00        10  
     2         0.75         0.80         0.78       130  
     3         0.68         0.73         0.70       132  
     4         0.64         0.55         0.59        42  
     5         0.00         0.00         0.00         5  
  
    accuracy          0.69  
   macro avg          0.34  
weighted avg          0.67
```